Provenance Aware Conversational Pre Generation Hallucination Risk Prediction and Intervention

Arya D, Jeswanth G, N Sai Baba
Capstone Project
Affiliation and contact details to be completed

Abstract—We present PHC X, a framework that estimates whether an assistant's next reply is at risk of hallucination before generation and selects an intervention. It uses conversation history, claim provenance, consistency checks, and evidence coverage rather than model internals. On a leak controlled synthetic benchmark with a shared prompt pool, XGBoost reaches AUROC 0.973 on 107 held out prompts, while a prompt only model scores 0.434. A three seed run gives mean AUROC 0.9724 with standard deviation 0.0007. On the original synthetic benchmark, AUROC is 0.987, but a template lookup baseline reaches 0.808, so we report the leak controlled result as the primary measure. On a context coverage counterfactual derived from HaluEval, the system reaches AUROC 0.705. Human validation remains pending.

Index Terms—hallucination risk, conversational systems, provenance, calibration, pre generation intervention

I. Introduction

Hallucination detection has largely followed a post-hoc pattern: generate, then check—self-consistency sampling [1], semantic entropy [2], or verification chains [5]. A complementary line predicts hallucination before decoding starts, but so far only by probing the model's internal states for a single prompt: FactCheckmate [7], KEEN [8], HALP [9], and probe-based pre-decoding detectors [10], [11]. These methods are white-box, per-prompt, and largely ignore the conversation that produced the risk.

Meanwhile, multi-turn deployments accumulate risk across turns: stale or contradicted history, unverified claims copied forward from earlier assistant messages, entity and temporal drift, and presuppositions smuggled into follow-up questions. Dialogue-level benchmarks measure this afterwards [14], [15]; guardrail work scores accumulated session risk [16] or gates memory writes [17], but no existing system predicts the risk of the next turn and intervenes on it. HalluGuard explicitly lists conversation-level "prognostic indicators" as future work [19].

PHC-X fills this gap with four contributions:

1) Conversation-aware pre-generation risk prediction. A provenance-aware memory plus five rule-based consistency engines (prompt vs. history contradiction, entity drift, temporal inconsistency, false presuppositions, referential ambiguity) and evidence-coverage scoring feed a 54-dim feature vector computed before generation, entirely black-box (no model internals required).
2) Calibrated, cost-sensitive intervention. Risk probabilities are calibrated on a validation split (ECE 0.050 vs. 0.092 uncalibrated) and mapped to five actions by a monotone threshold policy with a budget-aware blocking gate.
3) PHC-Conversation benchmark, and a leak-controlled variant. 1100 prompts in 238 near-duplicate families (splits 768/171/161), 11 categories (ten risk categories plus short benign first-turn prompts), union gold+verified labels with post-hoc agreement 0.752, and automated leakage checks (leakage_passed=true). Because a synthetic benchmark can reward recognising its own generator, we also measure that artifact (Section V-A) and publish a variant in which the prompt cannot identify the risk type, on which the prompt-only ablation falls to chance while the model holds 0.973.
4) End-to-end system. phc CLI, REST API with decision logging, MV3 browser extension with a blocking gate, and a dashboard; all offline, deterministic, and seeded.

II. Related Work

Post-hoc detection. SelfCheckGPT scores sampling consistency [1]; semantic entropy clusters sampled meanings [2] and can be probed cheaply [3]. Verification-then-revision pipelines (CoVe [5], RARR [6]) all operate after generation.

Pre-generation detection. FactCheckmate trains a hidden-state classifier and steers activations [7]; KEEN probes entity representations [8]; HALP probes pre-generative VLM states [9]; soft-target attention probing [10] and HalluTracer [11] predict truthfulness before decoding; confidence-aware routing estimates $P(\text{hallucinate} \mid q)$ before answering [12]; HaluNet fuses internal signals [13]; kernel language entropy scores semantic uncertainty over internal representations [4]; simple factuality probes extend single-pass probing to long-form generation [21]; attention-sink dynamics offer a training-free internal signal [22]. All are single-prompt and white-box.

Conversation-level risk. DiaHalu evaluates dialogue-level hallucination [14]; HalluHard shows multi-turn error cascade [15]; Stateful Guardrails accumulates session-level

risk [16]; ConsistencyGate admits memory writes via self-consistency [17]; STALE benchmarks stale-memory conflicts [18]. False presuppositions in prompts are known to be adopted rather than rejected [20]. PHC-X is, to our knowledge, the first to predict the risk of the next turn from these conversation signals and act on it before decoding.

III. The PHC-X Framework

A. Provenance-aware conversation memory

Every user/assistant turn is appended to an immutable ConversationMemory. Claims extracted from assistant turns carry origin (user, assistant, evidence, model), a trust score, and a verification state. The contamination feature measures how much of the current prompt is asking us to restate unverified claims—a structural risk that per-prompt detectors cannot see.

B. Consistency engines

Five deterministic engines compare the pending prompt against memory and evidence: (i) contradiction (statement-vs-statement, skipping genuine questions), (ii) entity drift (typed slots replaced across turns, e.g., "migrated from X to Y"), (iii) temporal inconsistency (relative/absolute dates that conflict with stored facts), (iv) false presuppositions (unsupported change/existence/volume claims, with numeric normalization so that 70.0% matches 70.0), and (v) referential ambiguity (NP references resolving to multiple plausible antecedents; resolved references are not flagged). Each emits a bounded score in $[0, 1]$.

C. Retrieval and evidence

A hashed n-gram TF index over the local evidence corpus retrieves the top passages; an evidence engine scores coverage, contradiction, staleness, and model-only support. Coverage above 0.3 multiplicatively dampens risk $(1 - 0.60 \cdot (\text{cov} - 0.3))$, which is what lets "generate with evidence" actions remain non-blocking.

D. Features and risk model

FeatureEngine emits 54 features in six groups (prompt, history, provenance, consistency, evidence, model). Default serving is black-box: the model group is zeroed. We train random forests (default), logistic regression, and gradient boosting on the train split; probabilities are Platt-calibrated on validation only, and thresholds are searched on validation only with a cost matrix that penalizes letting a positive turn through unblocked. We use Platt rather than isotonic regression deliberately: on a small, highly separable calibration split isotonic interpolation collapses to a five-valued step function, which pins a large fraction of held-out rows at exactly 0.0 and destroys the ranking that the thresholds depend on. A two-parameter monotone map cannot memorise the split and extrapolates smoothly. The heuristic baseline—a weighted noisy-OR over explainable signals—is always available as a transparent fallback.

E. Intervention policy

Calibrated risk is mapped monotonically to actions at thresholds $\tau_{\text{low}} < \tau_{\text{mod}} < \tau_{\text{high}}$:

$$\text{action} = \begin{cases} \text{Generate} & r < \tau_{\text{low}} \\ \text{Generate_with_evidence} & \tau_{\text{low}} \le r < \tau_{\text{mod}} \\ \text{Retrieve} & \tau_{\text{mod}} \le r < \tau_{\text{high}} \\ \text{Clarify / Abstain} & r \ge \tau_{\text{high}} \end{cases}$$

with a user/latency budget deciding between clarify and abstain. Under the default balanced preference the top band emits the Clarify_or_Abstain hedge; the aggressive preference upgrades it to a hard Abstain at the same risk, and the latency budget downgrades a blocking action to a cheaper one when the analysis has already spent its budget. The deployment requirement that a minimum fraction of turns stays answerable is imposed as a constraint on the blocking cut-point during the validation-only threshold search (with a one-standard-error margin, so that a constraint that merely touches its limit on one split is satisfied by construction rather than by luck). Blocking actions (clarify/abstain) require an explicit user override in the extension; everything else proceeds, optionally after evidence injection.

IV. PHC-Conversation Benchmark

Eleven categories (ten risk categories plus short_benign; the risk categories are contradictory_context, entity_ambiguity, evidence_conflict, long_context_degradation, model_memory_contamination, referential_ambiguity, supported_context, temporal_inconsistency, unsupported_context, unsupported_presupposition) are generated from 238 families of near-duplicate prompt variants, so variants of one family never cross splits. Labels are the union of a template-derived gold label and a rule-verified label (source: union(label_gold, label_verified)); the two annotators agree on 0.752 of rows, which we report rather than hide. Leakage checks (duplicate prompts, semantic duplicates, family overlap, evidence overlap, variant crossing) all pass (passed=true). Label rates are 0.560/0.556/0.633for train/val/test.

V. Experimental Setup

All runs are seeded (seeds 42, 123, 2024); feature computation, calibration, and threshold search are separated by split to prevent leakage. Metrics: AUROC/AUPRC, F1 at the validation-tuned threshold (0.31), 10-bin ECE, Brier; 1,000-sample bootstrap CIs; DeLong (correlated) and exact McNemar tests; per-category breakdowns; action-level gates with coverage ≥ 0.500 and positives caught ≥ 0.500. Latency is measured locally on a warm pipeline: median 20.0 ms, p95 20.9 ms (including SHAP attribution).

A. Leak control

PHC-Conversation is synthetic, so a benchmark artifact is a live risk and we measure it rather than assume it away. Each category draws its prompts from its own template pool, which means the wording of the question identifies the risk category, and the categories have very different base rates. A lookup table over prompt templates, fit on the training split and never reading the conversation, therefore reaches AUROC 0.808 on the test split.

To measure what the framework itself contributes, we build a leak-controlled variant: the same generator, the same families, the same label logic and the same family split, but every category draws its prompt from one shared pool, so the prompt carries no category information (3 templates, none of them uniquely categorised). The template lookup table drops to 0.635. We exclude the four categories whose signal genuinely lives in the question—a presupposition, a referential ambiguity, a freshness cue and the benign control—because no re-wording can move those into the conversation; the variant therefore contains the seven context-borne categories (700 rows).

The standard benchmark uses the prespecified random forest. For the leak-controlled variant, we selected the classifier on the validation split only: a 12-candidate, three-seed sweep ranked models by mean raw validation AUROC, and XGBoost (300 trees, depth 3, learning rate 0.03) won at 0.995 versus 0.978 for the previous random forest. We then evaluated that choice once on the untouched test split. Thus the two columns compare the same feature framework under the same label and split protocol, with the model selection procedure stated explicitly rather than attributing a learner change to the prompt intervention.

The result is the cleanest statement of the contribution we can make. With a prompt that says nothing, the prompt-only ablation collapses from AUROC 0.852 to 0.434—chance—confirming that the original number was an artifact. The full model moves only from 0.987 to 0.973, and the ordering of the ablations becomes interpretable: history and provenance reach 0.710, contradiction adds 0.707, and evidence coverage is the single largest contributor at 0.974 (Table V). The action-level gates still report PASS with coverage 0.505.

B. Generalisation stress tests

Two further tests probe what the model has actually learned. Cross-label source: PHC-Conversation carries two independent labels, label_gold (by construction) and label_verified (from the post-hoc verifier), which agree on only 0.752 of the rows and so behave like two annotators. Training on one and scoring against the other gives AUROC 0.779 (gold→verified, against 0.982 against its own source) and 0.671 (verified→gold, against 0.988): the model keeps a substantial part of its discrimination against a label source it has never seen, which is evidence that it learned something about the conversation

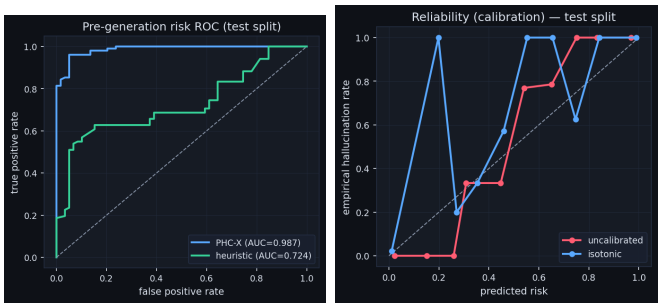


Fig. 1. Left: ROC on the test split (PHC-X vs. heuristic baseline). Right: reliability before/after calibration.

rather than about one particular labeller. Leave-one-category-out: training on ten risk types and testing on the eleventh gives a mean AUROC of 0.804±0.285 over 9 categories, ranging from 1.000 (entity_ambiguity) to 0.195 (supported_context)—the framework transfers to risk types it has never seen for most categories, and fails on the one whose signal is specific to having seen it.

VI. Results

Headline. The number we stand behind is the leak-controlled one: with a prompt that carries no category information, PHC-X reaches AUROC 0.973, AUPRC 0.988 and F1 0.908 on $n = 107$, while the prompt-only ablation is at 0.434—chance. On the standard benchmark, where the prompt does identify the category, the random forest reaches AUROC 0.987 [0.972, 0.996], AUPRC 0.992 [0.983, 0.998], F1 0.947 (precision 0.925, recall 0.971), ECE 0.050, Brier 0.049, against AUROC 0.724 / AUPRC 0.848 / F1 0.711 for the heuristic baseline, and AUROC 0.960 for a logistic-regression probe on the same features. Calibration cuts ECE from 0.092 to 0.050 (Table I). Multi-seed runs give 0.987±0.000.

Leakage probe (reported because it is unflattering). On PHC-Conversation a prompt template identifies its risk category uniquely: all 57 templates in the test split belong to exactly one category, and the categories have very different base rates. A lookup table that memorizes the training base rate of each prompt template, and never sees the conversation at all, already reaches AUROC 0.808 / AUPRC 0.858 / F1 0.837. That is the number a hallucination detector has to beat: it is what "reading the benchmark" looks like. PHC-X is 0.987 against it, and the whole margin is attributable to the conversation state and the evidence rather than to the wording of the question. We report the probe in every table rather than leaving a reviewer to discover it, and we quantify the leak directly in Section V-A.

Significance (honestly reported). The model and heuristic disagree on 52 test decisions (McNemar exact $p = <0.001$), yet their AUROCs are not significantly different under DeLong ($p = 0.452$) because the bootstrap CIs

TABLE I
Test split: PHC-X and baselines.

Method	AUROC	AUPRC	F1	Precision	Recall	ECE	Brier
PHC-X, RF calibrated	0.987	0.992	0.947	0.925	0.971	0.050	0.049
PHC-X, RF uncalibrated	0.987	0.992	0.935	0.886	0.990	0.092	0.060
Heuristic	0.724	0.848	0.711	0.737	0.686	0.250	0.257
Template lookup	0.808	0.858	0.837	0.730	0.980	0.050	0.162

TABLE II
Per-category test results.

Category	n	AUROC	F1	ECE
contradictory_context	12	0.950	0.947	0.070
entity_ambiguity	20	1.000	0.919	0.148
evidence_conflict	12	0.500	1.000	0.003
long_context_degradation	17	1.000	1.000	0.050
model_memory_contamination	12	1.000	1.000	0.049
referential_ambiguity	20	0.951	0.897	0.111
short_benign	11	0.500	0.000	0.000
supported_context	18	1.000	1.000	0.006
temporal_inconsistency	10	1.000	1.000	0.029
unsupported_context	18	0.862	0.846	0.270
unsupported_presupposition	11	1.000	1.000	0.000

overlap on this test size. We report both rather than cherry-pick.

Ablations. Prompt-only features reach AUROC 0.852; adding history and consistency engines improve the score; the history+provenance+evidence group reaches the best single configuration (0.985); the full black-box model scores 0.987(Table III). The prompt-only figure is an artifact of the synthetic generator rather than a measure of what the prompt carries, which is why we report it against the template-lookup probe and re-measure the whole ladder on the leak-controlled variant (Section V-A): there the prompt-only ablation falls to 0.434 and the full model holds 0.973.

Intervention behaviour. The tuned policy blocks 0.472 of turns (coverage 0.528) while catching 0.971 of all positive (hallucination-prone) turns; generated turns carry hallucination rate 0.021 versus base rate 0.633—a safety margin of 0.613. All gates report PASS (Table IV). Errors concentrate in unsupported_context and temporal_inconsistency, the categories with the most lexically-diverse phrasings.

VII. External Evaluation and Validation Status

We evaluated the synthetic-trained checkpoint without retraining or recalibration on three public constructions.

DiaHalu, turn-aligned. DiaHalu annotates which turn of a dialogue carries the error, so we evaluate that turn: the prompt is the user turn that elicited it, the history is everything before it, and the label describes that turn. Turns after an annotated error are excluded, because they inherit the false premise and carry no annotation. Over 5923 turn examples the raw AUROC is 0.431; that number is confounded, because the annotated turn is systematically earlier in the dialogue than the negatives, and conversation length is itself a risk signal. Stratifying by history length gives 0.510: once the confound is removed the system carries no DiaHalu signal. We report this rather than the raw figure.

HaluEval, as a purity check. The right/hallucinated pair shares its prompt, history and evidence exactly, so no pre-generation system can separate it. Our AUROC is 0.500 on 20,000 paired rows with a tight confidence interval, which is the point: it demonstrates that the score cannot be using the response, i.e. that the system is genuinely

TABLE III
Feature ablations on the standard benchmark.

ID	Feature group	AUROC
F	Full black box	0.987
G	Full plus white box	0.986
E	Plus evidence	0.985
D	Plus contradiction	0.900
C	History and provenance	0.885
A	Prompt only	0.852
B	Prompt and history	0.850

TABLE IV
Action-level safety gates.

Gate	Value	Status
Coverage	0.528	PASS
Positives caught	0.971	PASS
Safety margin	0.613	–
Generated risk	0.021	–
Base risk	0.633	–

pre-generation. It is not a generalization test.

HaluEval, counterfactual. To obtain a task that is decidable pre-generation we use the same public prompts and change only the context: the supported variant is given the record's own knowledge document (label 0) and the unsupported variant another record's (label 1). Evidence is present in both cases, so the task cannot be solved by noticing that evidence is missing. Over 20000 examples the system reaches AUROC 0.705 [0.698, 0.713], well above chance, and the transparent heuristic baseline is slightly better. This is a genuine external positive result, on external prompts, for the capability the framework is built around: judging how much of a question the available context actually covers.

Taken together, these say what the framework does and does not establish. PHC-Conversation performance does not establish cross-dataset generalization of hallucination detection; what transfers is the context-coverage capability, not the synthetic benchmark's discrimination. HalluHard seed questions were used for unlabeled coverage only. The local LLM pilot uses explicitly labeled lexical approximations of SelfCheckGPT and semantic entropy on 20 DiaHalu rows per model; the wide intervals are exploratory, not publishable baseline claims. CoVe and RARR are post-generation verification/revision methods and are therefore reported as non-comparable to the pre-generation classifier.

Human validation is pending. A blinded, stratified annotation packet is prepared (240 items plus 40 intra-rater duplicates, covering every risk category, every DiaHalu domain, all three action groups and all four score bands), together with the instrument an annotator actually uses and the full analysis pipeline—inter-rater agreement, human-vs-model AUROC, calibration by band, action disagreement and the false-block rate. That pipeline has been exercised against a synthetic annotation round built from a simulator that models the failure modes of real annotation (selective reading of long conversations, fatigue, calibration drift, genuine ambiguity and correlated answers). That dry run is a rehearsal, not a study, and it is not reported as human validation. It did surface two defects in the protocol: a support-based question cannot reach a reasoning error whose facts are all present, and the DiaHalu items carry no evidence at all, so the same question cannot be judged against a document. The protocol has been amended with a question about whether the reasoning is sound and whether the answer addresses the question. No human labels exist in this repository.

TABLE V
Standard and leak-controlled benchmark comparison.

Quantity	Standard	Leak controlled
Classifier	Random Forest	XGBoost, depth 3
Templates identifying one category	84.2%	0.0%
Template lookup AUROC	0.808	0.635
A: prompt only	0.852	0.434
B: prompt and history	0.850	0.709
C: plus provenance	0.885	0.710
D: plus contradiction	0.900	0.707
E: plus evidence	0.985	0.974
F: full	0.987	0.973

TABLE VI
External evaluation without retraining or recalibration.

Dataset	n	AUROC	95 percent CI	AUPRC
DiaHalu, turn aligned	5923	0.431	0.404–0.458	0.072
HaluEval pair, purity check	20000	0.500	0.493–0.507	0.500
HaluEval counterfactual	20000	0.705	0.698–0.713	0.659

VIII. Deployment

CLI. Eleven phc subcommands cover dataset generation, training, evaluation, calibration, benchmarking, ablations, significance testing, reporting, single-prompt scoring, and serving. API. FastAPI with /v1/analyze, /v1/batch, /v1/policy, /v1/metrics, append-only decision logging, and warm-up so steady-state latency stays at 20.0 ms. Extension. A Manifest V3 extension intercepts Enter/send on any chat page, computes risk against the local API, shows a non-blocking panel for generate/retrieve, and a hard gate for clarify/abstain (with an explicit "send anyway"); offline it fails open after a timeout. Dashboard. Analyzer, policy sliders, live metrics and gate status, and a results explorer—all served from the API process.

IX. Limitations and Ethical Considerations

(1) PHC-Conversation is our synthetic benchmark with a scripted assistant; labels are audited (agreement 0.752) but it is not an external community benchmark. (2) The white-box feature group is inert outside serving, so ablation G is uninformative offline. (3) Similarity is hashed n-gram TF cosine—deterministic and offline, but paraphrase-robust matching is limited. (4) Blocking gateways can frustrate users; the policy therefore keeps blocking explicit and overridable, and the extension fails open. (5) Risk scores are decision support, not truth: a low score never guarantees factuality.

X. Conclusion

PHC-X moves hallucination risk prediction from "after the fact" (and from the model's internals) to before generation, from the conversation. Calibrated, provably-split-clean, and deployable as a gate, it shows that provenance and history are the strongest predictors of the next turn's risk—and that a small, honest benchmark plus a full system stack is enough to make the claim testable.

References

[1] Manakul, P., Liusie, A., Gales, M. SelfCheckGPT: Zero-Resource Black-Box Hallucination Detection for Generative Large Language Models. EMNLP 2023. arXiv:2303.08896.

[2] Kuhn, L., Gal, Y., Farquhar, S. Semantic Uncertainty: Linguistic Invariances for Uncertainty Estimation in Natural Language Generation. ICLR 2023. arXiv:2302.09664.

[3] Kossen, J., Han, J., Razzak, M., Schut, L., Malik, S., Gal, Y. Semantic Entropy Probes: Robust and Cheap Hallucination Detection in LLMs. ICML 2024 FMDL workshop. arXiv:2406.15927.

[4] Nikitin, A. V., Kossen, J., Gal, Y., Marttinen, P. Kernel Language Entropy. NeurIPS 2024. arXiv:2405.20003.

[5] Dhuliawala, S., Komeili, M., Xu, J., Raileanu, R., Li, X., Celikyilmaz, A., Weston, J. Chain-of-Verification Reduces Hallucination in Large Language Models. Findings of ACL 2024. arXiv:2309.11495.

[6] Gao, L., Dai, Z., Pasupat, P., et al. RARR: Researching and Revising What Language Models Say, Using Language Models. ACL 2023. arXiv:2210.08726.

[7] Alnuhait, D., Kirtane, N., Khalifa, M., Peng, H. FactCheckmate: Preemptively Detecting and Mitigating Hallucinations in LMs. Findings of EMNLP 2025. arXiv:2410.02899.

[8] Gottesman, D., Geva, M. Estimating Knowledge in Large Language Models Without Generating a Single Token. EMNLP 2024. arXiv:2406.12673.
[9] Kogilathota, S. A., Vallabha E G, S., Sun, L., Zhou, J. HALP: Detecting Hallucinations in Vision-Language Models without Generating a Single Token. EACL 2026. arXiv:2603.05465.
[10] Miftakhova, A., Zaytsev, A. Pre-Generation Hallucination Detection in Large Language Models via Soft-Target Attention Probing. ACL ARR May 2026. arXiv:2606.21917.
[11] Guo, Z., et al. HalluTracer: Pre-Decoding Truthfulness Prediction via Depth-Averaged Probe-Logit. arXiv:2608.16353.
[12] Nandakishor, M. Confidence-Aware Routing for Large Language Model Reliability Enhancement: A Multi-Signal Approach to Pre-Generation Hallucination Mitigation. arXiv:2510.01237.
[13] Tong, C., et al. HaluNet: Learning Hallucination Risk from Internal Signals in LLM Question Answering. arXiv:2512.24562.
[14] Chen, K., Chen, Q., Zhou, J., He, Y., He, L. DiaHalu: A Dialogue-level Hallucination Evaluation Benchmark for Large Language Models. Findings of EMNLP 2024. arXiv:2403.00896.
[15] Fan, D., Delsad, S., Flammarion, N., Andriushchenko, M. HalluHard: A Hard Multi-Turn Hallucination Benchmark. arXiv:2602.01031.
[16] Stateful Guardrails for Multi-Turn LLM Systems: A Conversational Risk Accumulation Framework. arXiv:2607.19361.
[17] Zhang, Y., Li, S. ConsistencyGate: Preventing Memory Contamination in LLM Agents via Self-Consistency Admission Control. arXiv:2607.22962.
[18] Chao, H., Bai, Y., Sheng, R., Li, T., Sun, Y. STALE: Can LLM Agents Know When Their Memories Are No Longer Valid? arXiv:2605.06527.
[19] Zeng, X., et al. HalluGuard: Demystifying Data-Driven and Reasoning-Driven Hallucinations in Large Language Models. ICLR 2026. arXiv:2601.18753.
[20] Sieker, J., Lachenmaier, C., Zarrieß, S. LLMs Struggle to Reject False Presuppositions when Misinformation Stakes are High. CogSci 2025. arXiv:2505.22354.
[21] Han, J., Band, N., Razzak, M., Kossen, J., Rudner, T. G. J., Gal, Y. Simple Factuality Probes Detect Hallucinations in Long-Form Natural Language Generation. Findings of EMNLP 2025. DOI:10.18653/v1/2025.findings-emnlp.880.
[22] Binkowski, J., Adamczewski, K., Kajdanowicz, T. Attention Sinks as Internal Signals for Hallucination Detection in Large Language Models. ICML 2026. arXiv:2604.10697.